



UNIVERSITÀ
DI TRENTO

Deliveroo.js

Autonomous Software Agents

a.y. 2025/2026

Davide Donà - 264467

Andrea Blushi - 264468

June 2026

Contents

1 Problem Statement	1
2 BDI Agent	1
2.1 Belief Revision	2
2.1.1 Agent beliefs	2
2.1.2 Map beliefs	2
2.1.3 Parcel beliefs	2
2.2 Desires and Intention Selection	3
2.2.1 Queue reconstruction	3
2.2.2 Dynamic Scoring	3
2.3 Planning: A* Navigation and Collision Handling	5
3 PDDL-Based Crate Navigation	5
3.1 Domain Overview	5
3.2 Planning With PDDL	6
3.3 A Smart Solution — Problem Restatement & Caching	7
4 LLM-Based Agent Layer	7
4.1 Architecture and Request Handling	7
4.2 Adaptive Behaviour	8
5 Coordination Strategies	8

1 Problem Statement

The `Deliveroo.js` environment is a grid-based delivery simulation where autonomous agents collect and deliver scattered parcels to designated tiles to maximise their accumulated score. The environment introduces various challenges: agents only perceive tiles within a limited sensor radius at any given time step. Parcel rewards decay at a fixed server-configured interval, creating time pressure that incentivises early collection. Multiple agent types may share the map; agents belonging to the same team cooperate to jointly maximise reward, whereas agents from different teams compete for the same parcels. Additionally, movable crates often obstruct navigable paths, meaning agents must actively push them to clear blocked routes.

We implement a BDI (Belief-Desire-Intention) agent augmented with an optional Large Language Model (LLM) based control layer. The BDI core handles reactive parcel collection, exploration, and crate-clearing autonomously; the LLM layer intercepts high-level instructions, adapts agent behaviour through a set of structured tools, and coordinates multi-agent strategy via periodic team assignments.

2 BDI Agent

The agent operates on a continuous cycle driven by the `sensing` socket event emitted by the server. Rather than relying on a static script, the agent dynamically reacts to the environment through a strict, four-stage pipeline. Upon receiving a `sensing` event, the following sequence executes:

1. **Perception** (`perceive()`): The agent integrates the new sensory data into its belief state through a Belief Revision process (detailed in Section 2.1). The actions we perform that modify the belief state automatically updates the beliefs, without having to wait for the next sensing event.
2. **Deliberation** (`deliberate()`): Using its updated beliefs, the agent generates potential goals (desires) and calculates a score for each. It then rebuilds its intention queue, prioritizing the most valuable task (detailed in Section 2.2).
3. **Planning** (`Planner.plan()`): The planner extracts the highest-ranked intention from the queue and translates it into a concrete sequence of executable, typed steps (detailed in Section 2.3).
4. **Execution** (`Executor`): Finally, the planned actions are asynchronously emitted to the server via socket events. To maintain responsiveness, the executor triggers an action at each server clock tick, instead of waiting for a perception update or an acknowledgment from the server.

2.1 Belief Revision

The agent’s belief state relies on two foundational utility data structures to manage temporal observations, which in turn feed into three specialized sub-systems: **AgentBeliefs**, **MapBeliefs**, **ParcelBeliefs**.

Tracker **Tracker** stores a single latest observation $\langle v, t_{seen} \rangle$ per entity **e**. On each sensing report, **update(e, v)** overwrites the existing entry. The confidence of a tracked observation decays exponentially over time:

$$c(t) = \exp\left(-\frac{t - t_{seen}}{\tau}\right) \quad (1)$$

where τ is the exponential time constant (the duration after which confidence drops significantly).

When the current sensing window covers an entity’s last-known position, but the entity is absent from the report, **invalidateAtSensedPositions** nullifies that position without removing the record entirely.

Memory While the tracker maintains only the latest state, **Memory** is a supplementary store that maintains a bounded temporal history of observations per entity, automatically evicting entries older than a defined time threshold.

2.1.1 Agent beliefs

AgentBeliefs stores the latest state of every observed agent via **Tracker**, and a bounded position history for each enemy via **Memory**. It also maintains a spatial heat field over the map, aggregating enemy presence across all tracked observations.

2.1.2 Map beliefs

MapBeliefs stores the static tile layout, a set of transient blocked tiles each with a configurable TTL. Upon map receipt, it performs connectivity analysis to seal *sink tiles*, walkable tiles unreachable from any delivery point, reclassifying them as walls.

2.1.3 Parcel beliefs

This component tracks the availability and value of delivery targets. Parcels within the current sensing window receive their authoritative reward from the server. Once out of sight, parcel rewards decay linearly at -1 per server-defined interval Δ_t . Each parcel maintains an independent decay clock to prevent drift accumulation. Given a parcel last observed at time t_0 with reward r_0 , its estimated reward at time t is:

$$r(t) = r_0 - \left\lfloor \frac{t - t_0}{\Delta_t} \right\rfloor. \quad (2)$$

When $r(t) \leq 0$, the parcel is considered expired and is deleted from the agent's beliefs.

2.2 Desires and Intention Selection

For the BDI layer, we define a fixed set of 3 core desire types, each parameterized by a target tile:

- **ReachParcel**: navigate to and collect a visible parcel;
- **DeliverParcel**: navigate to a delivery tile while carrying parcels;
- **Explore**: navigate to an unobserved spawn tile when idle (i.e., no reachable parcels or deliveries).

2.2.1 Queue reconstruction

At each deliberation step, the intention queue is fully reconstructed from the current belief state. The queue is sorted by priority first, then by descending score within each desire type. Both **ReachParcel** and **DeliverParcel** desires have a priority of 1, while **Explore** has a lower priority of 0. This ensures that the agent always prefers tangible parcel-related goals over exploration when any are available.

2.2.2 Dynamic Scoring

The agent computes a dynamic score for each desire to guide intention selection. The score is a real-valued estimate of the desire's expected utility, incorporating factors such as the parcel's reward, the distance to the target, and the presence of competitors.

Distance Penalty Each desire's score is inversely proportional to the estimated distance to its target, calculated using the agent's current beliefs. The agent computes the shortest path distance d from its current position to all other tiles using a Breadth-First Search (BFS). We will refer to this distance as d throughout the rest of the section. (Note: when $d = 0$, we assign an infinite score).

ReachParcel Let r_p denote the current estimated reward of a target parcel p , and C the set of parcels the agent is currently carrying. The score S for a **ReachParcel** desire targeting parcel p is calculated as:

$$S = \frac{r_p + \sum_{c \in C} r_c}{d} \cdot \gamma$$

Where γ is a race factor that de-prioritizes parcels likely to be collected by competitors. It is computed as:

$$\gamma = \min\left(1, \frac{d_{\text{enemy}}}{d}\right)$$

where d_{enemy} is the distance from the nearest tracked enemy to the parcel.

Incorporating the total reward of currently carried parcels ($\sum_{c \in C} r_c$) into the score allows the agent to evaluate the utility of picking up a new parcel against the potential value of delivering existing ones. This effectively balances the agent's acquisition and delivery goals.

DeliverParcel For a **DeliverParcel** desire, the objective is to secure the rewards of the parcels currently held by the agent. The score S is computed as the total estimated reward of all carried parcels C , inversely scaled by the shortest path distance d to the delivery tile:

$$S = \frac{\sum_{c \in C} r_c}{d}$$

Explore The **Explore** desire serves as a fallback to locate newly spawned parcels. We first define the normalized observation age τ :

$$\tau = \min \left(1, \frac{\Delta t}{T_{\text{spawn}}} \right)$$

where Δt is the time since the target spawn tile x was last observed, and T_{spawn} is the environment's spawn interval.

The cluster weight w quantifies the structural density of the target's surrounding spawn area. It is computed as the sum of inverse distances to all neighbor spawn tiles within the observation range R_{obs} :

$$w = \sum_{y \in OST} \frac{1}{\delta(x, y) + 1}$$

where OST is the set of spawn tiles such that their Manhattan distance $\delta(x, y) \leq R_{\text{obs}}$.

The enemy heat h actively steers the agent away from sectors currently crowded by competitors. It is accumulated from the history of tracked enemy observations E , combining spatial and temporal decay:

$$h = \sum_{e \in E} \exp \left(-\frac{d_e^2}{2\sigma^2} \right) \exp \left(-\frac{\Delta t_e}{\tau_h} \right)$$

where d_e is the spatial distance to the past enemy observation, Δt_e is the observation's age, σ controls the spatial spread, and τ_h is the temporal decay rate.

Finally, the score for an **Explore** desire is computed as:

$$S = \frac{w \cdot \tau}{d \cdot (1 + h)}$$

2.3 Planning: A* Navigation and Collision Handling

For pathfinding, the agent employs a standard A* search algorithm with a Manhattan distance heuristic, which is admissible in our grid-based environment. The planner incorporates dynamic walkability checks based on the agent’s current beliefs, treating temporarily blocked tiles as impassable. Statically, walls and out-of-bounds tiles are always rejected; conveyor tiles are additionally direction-sensitive, blocking movement against their flow.

Plans generated by the A* algorithm are defined as sequences of typed actions: directional **move** steps followed by a terminal **pickup** or **putdown** step. An existing plan is reused across deliberation steps as long as the active desire type and target remain unchanged, and all remaining steps successfully pass walkability re-validation.

Collision Handling Prior to executing a move step, the agent verifies the availability of the target tile. A tile is considered occupied, and thus temporarily blocked, under two conditions: either an enemy is currently observed on it, or an enemy’s predicted next position matches it with sufficient confidence.

When the next planned tile is occupied, the collision manager applies a three-stage escalation:

1. **Detour:** replan with the blocked tile treated as a wall. This is accepted only if the detour adds at most a predefined threshold of Δ_{th} steps over the remaining plan.
2. **Wait:** if no acceptable detour exists, pause execution for a random duration drawn from a predefined uniform distribution.
3. **Block and replan:** after the wait expires, or after a specified limit L of repeated invalidations, mark the tile as temporarily blocked for a randomized time-to-live (TTL) duration and completely replan toward the same goal.

3 PDDL-Based Crate Navigation

Some maps in the Deliveroo.js environment contain movable crates that can obstruct navigable paths. An agent may push a crate one tile in any cardinal direction, subject to the constraints outlined in Section 2.3. We model this crate manipulation problem using PDDL (Planning Domain Definition Language), invoking it as a fallback mechanism when the standard A* planner fails to find an unobstructed route.

3.1 Domain Overview

The `deliveroo-crates` domain is specified in STRIPS, utilizing two core object types: `tile` and `crate`. The predicate set captures the complete state required for crate-aware navigation:

- (**at** ?*t*): the agent occupies tile *t*;
- (**adj-up/down/left/right** ?*from* ?*to*): defines the directed spatial relationship between adjacent tiles, specifying valid movement or pushing directions;
- (**crate-at** ?*c* ?*t*): crate *c* is positioned on tile *t*;
- (**crate-free** ?*t*): tile *t* currently contains no crate, meaning it is available for agent entry or as a push destination;
- (**crate-space** ?*t*): tile *t* is structurally capable of supporting a crate (e.g., it is not a wall or a conveyor belt).

The domain defines eight parameterized actions: four directional *move* actions and four directional *push* actions. A move action allows the agent to step into an adjacent **crate-free** tile. A push action allows the agent to displace a crate one tile in the specified direction, requiring two preconditions: (i) the agent is positioned adjacent to the crate on the opposite side of the push direction, and (ii) the target destination is a valid **crate-space** and is currently **crate-free**. The execution of a push action simultaneously relocates the crate to the target tile and moves the agent onto the crate’s former tile, updating the **crate-free** state of the involved tiles accordingly.

3.2 Planning With PDDL

PDDL planning is triggered exclusively when A* fails and the belief system confirms that crates obstruct the route (Section 2.3). The problem instance is generated by **buildProblem**: it emits adjacency facts for every non-wall tile pair, **crate-at** facts for each believed crate, **crate-free** facts for all unoccupied tiles, and **crate-space** facts for all structurally valid positions. The goal is a single ground atom (**at target**).

Two invocation modes exist. In the *stitched* mode, the planner receives only the minimal entry–exit cluster segment identified by **findCrateSegment** (Section 2.1); the A* prefix to the cluster entry and the A* suffix from the cluster exit to the final goal are computed separately and concatenated. In the *whole-trip* fallback mode, the full route from the agent’s current position to the goal is submitted when the segmentation heuristic fails (e.g. multi-cluster corridors).

Solver selection is runtime-configurable: when the environment variable **PAAS_HOST** is set, problems are submitted to the online **@unitn-asa/pddl-client** solver; otherwise the local binary **dual-bfws-ffparser** (accessed via **planutils**) is invoked in a per-call temporary directory to avoid file-system races in multi-agent deployments. Both **move-*** and **push-*** steps in the returned plan are translated into a **move** plan step (the agent’s tile after the action), and appended to the A* prefix.

3.3 A Smart Solution — Problem Restatement & Caching

Since the PDDL planner introduced overhead that was not acceptable for the reactive nature of the environment, we implemented a problem restatement to enable plan caching.

Instead of invoking the planner for a global solution from the agent’s current position to the final goal, the plan is decomposed into three steps:

- Move from the current position to the cluster entry tile (A^*);
- Navigate through the crate cluster (PDDL);
- Move from the cluster exit tile to the final goal (A^*).

The PDDL planner is therefore restricted to the second step alone. Computed maneuvers are stored in a FIFO cache of size 32, keyed by the tuple (entry, exit, {crate positions}). The cache key is independent of the agent’s absolute position: it depends only on the relative crate arrangement and the cluster’s entry and exit tiles. Consequently, the exact same maneuver plan is reused whenever the agent encounters a structurally identical sub-problem — including cases where the agent approaches the same cluster from a different direction — drastically reducing planner invocations and keeping the agent reactive.

4 LLM-Based Agent Layer

4.1 Architecture and Request Handling

The LLM layer is implemented as a thin wrapper around the BDI agent. `LLMAgent` instantiates a `BDIAgent` and an `LLMClient`, wiring the latter with callbacks into the BDI core: `addInjectedIntention`, `removeIntentionsByType`, a shared `Messenger`, and a shared `RuleStore`.

Incoming messages arrive via the `msg` socket event. The handler applies two filters in sequence: (i) coordination messages with a `beliefs_report` envelope are forwarded to the `Coordinator` regardless of sender; (ii) all other messages are accepted only from the sender specified by the environment variable `MISSION_AGENT_NAME`, acting as the mission controller. Valid messages are passed to `LLMClient.processMessage`.

Tool execution follows a multi-hop loop. Each invocation calls the OpenAI-compatible API with `tool_choice`: "auto" and `temperature`: 0.1 for up to `maxHops` = 5 rounds. Tool results are appended as `role`: "tool" messages and the next hop proceeds. The loop continues automatically only for the `calculate` tool (arithmetic evaluation), whose result is typically consumed in a follow-up reasoning step; all other tool calls terminate the loop after execution.

4.2 Adaptive Behaviour

The LLM is exposed to eleven tools that allow it to modify the agent’s behaviour at runtime without stopping execution:

- **request_goto / assign_goto**: inject a `REACH_TILE` desire with a configurable reward and expiry time; **assign_goto** addresses a specific agent by id, sending the instruction peer-to-peer when targeting a teammate.
- **request_putdown_at**: inject a `DELIVER_PARCEL` desire targeting a specific tile, with an LLM-specified bonus reward to bias delivery scoring.
- **register_scoring_rule**: upsert a scoring rule (stack-count, delivery-tile, or parcel-value axis) into the shared `RuleStore`, directly modifying the desire ranking formulas described in Section 2.2.
- **register_traversal_penalty**: register a soft cost $\delta(p)$ on a tile, affecting both A* routing and penalized distance scoring (Equation (??)).
- **request_rendezvous**: inject a `HOLD_TILE` desire with a release zone around the target; both agents independently converge to the joint-optimal holding tile and auto-release once the zone condition is satisfied.
- **request_red_light / request_green_light**: inject an indefinite `HOLD_TILE` on the nearest odd-row tile (red light), or clear all `HOLD_TILE` desires (green light), halting and resuming the full team.
- **request_team_status**: trigger an immediate coordination round (Section 5).
- **reply**: send a text message back to the mission controller via the socket.
- **calculate**: evaluate a sandboxed arithmetic expression (the only auto-continuing tool).

All tools that affect agent behaviour share a common injection path (`applyInjection`) with the peer-message dispatcher, ensuring that LLM-issued commands and peer-forwarded commands follow identical validation and application logic.

5 Coordination Strategies

Peer communication. The `Messenger` wraps the socket’s `emitSay` and `emitShout` primitives with a friend-gating filter: only messages from agents on the same team are dispatched. Outgoing tool invocations are encoded as peer-injection envelopes (format `{v:1, kind:"peer_injection", tool, args}`) and broadcast to all friends via `communicate`. Upon receipt, each teammate decodes and applies the same injection locally, so a single LLM tool call propagates atomically to the whole team. Point-to-point delivery via `communicateTo` is used when a command targets a specific agent.

Rendezvous. When the LLM issues `request_rendezvous(x, y, max_distance)`, all candidate tiles within Manhattan distance `max_distance` of the target are injected as `HOLD_TILE` desires with a release zone centred on (x, y) . Both agents independently score these tiles using the bottleneck rendezvous formula (Sec-

tion 2.2), selecting the same joint-optimal tile without explicit negotiation. The hold intention is automatically pruned once all agents are inside the release zone.

Red-light / green-light. The red-light mechanism injects an indefinite `HOLD_TILE` desire on the nearest odd-row tile for each agent. Because `HOLD_TILE` holds priority 1, it immediately supersedes all delivery and parcel desires. A safety expiry TTL is attached to prevent agents from being stranded if the green-light signal is lost. The green-light command clears all `HOLD_TILE` desires on both agents, resuming normal operation. Both commands are broadcast to teammates, so the entire team stops and resumes together.

Coordinator. A periodic `Coordinator` runs a team assignment round every 10 000 ms (with a 5 000 ms cooldown to prevent re-entrancy). Each round: (i) broadcasts a `request_beliefs` envelope to all friends; (ii) waits 750 ms for belief reports to arrive; (iii) filters reports to those fresher than the request timestamp; (iv) builds a geometric summary of the team state (*hot zones*, agent positions, carried parcels); (v) passes this summary to `LLMClient.coordinate`, which issues `assign_goto` calls to direct each agent towards a high-value zone. The `request_team_status` tool allows the mission controller to trigger this round on demand.

Cooperative vs. competitive behaviour. The friend/enemy distinction is determined solely by team membership. Communication, peer injection, and coordination are restricted to friends. Enemy agents contribute to three belief-level mechanisms: (i) the race factor γ in `REACH_PARCEL` scoring, which discounts parcels reachable earlier by a competitor; (ii) the enemy heat field H (Equation (??)), which suppresses exploration towards contested zones; (iii) trajectory prediction and path blocking, which prevents the agent from planning moves into tiles likely occupied by an enemy. No messages or injections are ever accepted from or sent to enemy agents.